

Machine Learning Engineer Nanodegree

Capstone Project Report

Felix H.L Lou
July 6th, 2018

Multi-channel Convolutional Neural Network for Text Classification

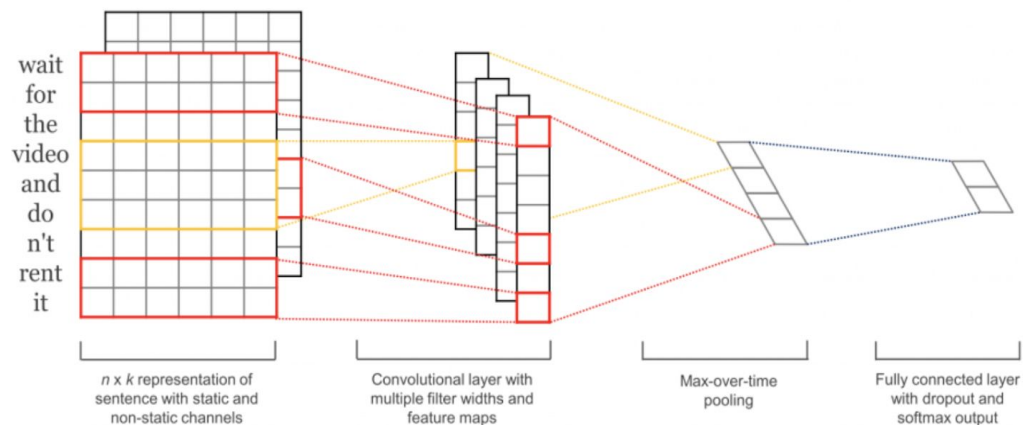
1. Definition

Problem Overview

Natural Language Processing, a.k.a NLP, is a study of automatic manipulation of natural language. In other words, it focuses on the interactions between human languages and computers, specifically on how to program computers to analyze and process a huge amount of natural language data.

Deep Learning is one of the members of the Machine Learning family. Originally inspired by how the human brain functions, Deep Learning models can be enhanced by training them with more examples whilst increasing their depth (Layers), so-called Deep Neural Network.

Businesses have been getting progressively more interested in applying these tools to business analytics. Sentiment Analysis is perhaps the most common application where NLP and Machine Learning shine. Sentiment essentially relates to feelings, that is, attitudes, emotions, or opinions. With Machine Learning, businesses leverage computational power to perform text analytics in a much more scalable way, in which they get to understand the attitudes of customers with respect to certain products/ services based on feedback. In other words, they are able to tell polarity (Positive/ Neutral/ Negative). In this project, we leveraged Convolutional Neural Networks with different kernel (gram) sizes to read customers' reviews and perform classification.



Problem Statement

a. Description

In this project, we leveraged Deep Neural Network to perform supervised learning, in which the model learned to ‘read’ customers’ reviews to tell the relevant sentiment. In other words, the goal was to build a model to detect customers’ emotional tone in their reviews. This is a typical Sentiment Analysis, in the form of text classification problem. Social media is perhaps the one that encounters this problem the most. As an example, the practitioners from Stanford University trained a classifier with supervised learning to tell people’s sentiment on Twitter:

<https://cs.stanford.edu/people/alecmgo/papers/TwitterDistantSupervision09.pdf>

b. Breakdown

- i. Task: Classify reviews’ potential sentiment (Emotional tone)
- ii. Performance: Accuracy/ F-score of the predictions on a test set
- iii. Independent variable (Predictor): Customers’ review
- iv. Dependent variable (Outcome): Customers’ sentiment (Positive/ Negative/ Neutral)

Metrics

In a supervised learning setting, observations are already labeled based on ground truth, as if there was an oracle. Using accuracy as a metric for evaluating a particular model's performance would therefore be appropriate. Nonetheless, accuracy is greatly affected by the distribution of classes. For instance, if most customers have negative reviews, we could say a particular customer is sad/ mad and could still be generally right, which is in fact missing the point of building a Machine Learning model. This is where precision and recall kicks in, and a common metric for classification problems is f-score, which considers both precision and recall to evaluate models’ performance:

$$F_{\beta} = (1 + \beta^2) \cdot \frac{\text{precision} \cdot \text{recall}}{(\beta^2 \cdot \text{precision}) + \text{recall}}$$

Whether to lean towards precision or recall depends on businesses. For instance, if identifying a negative user review as a positive one would mean fatel to the business, then it is probably way more important for the model to precisely predict negative reviews than to recall them. In this project, we used f-1 score, which captures the harmonic mean of precision and recall to objectively evaluate the model’s performance:

$$\text{F1 Score} = 2 \cdot \frac{\text{Precision} \cdot \text{Recall}}{\text{Precision} + \text{Recall}}$$

2. Analysis


Data Exploration

The dataset originally included 23486 observations and 10 independent variables. Each observation corresponds to a customer's review, and includes other variables. Additional columns like 'Word Count' and 'Customer Satisfaction' were created for exploratory purposes and statistical analysis. Columns regarding sentiments were also created with the help of the NLTK package. There are about 844 observations missing certain values, which were dropped when preprocessing; so the actual observations that later fed into the model was only 22628. The index column in the original csv file was also dropped.

```
Data columns (total 17 columns):
Clothing ID          22628 non-null int64
Age                  22628 non-null int64
Title                19662 non-null object
Review Text          22628 non-null object
Rating               22628 non-null int64
Recommended IND      22628 non-null int64
Positive Feedback Count 22628 non-null int64
Division Name        22628 non-null object
Department Name      22628 non-null object
Class Name           22628 non-null object
Word Count           22628 non-null int64
Customer Satisfaction 22628 non-null int64
Polarity Score        22628 non-null float64
Positive Score        22628 non-null float64
Neutral Score         22628 non-null float64
Negative Score        22628 non-null float64
Sentiment             22628 non-null object
```

'Word Count' is based on 'Review Text'. And in order to get 'Customer Satisfaction', we leveraged the mean of 'Rating' to get a threshold of around 4. This was proved to be probabilistic through bootstrapping (10000 simulations):

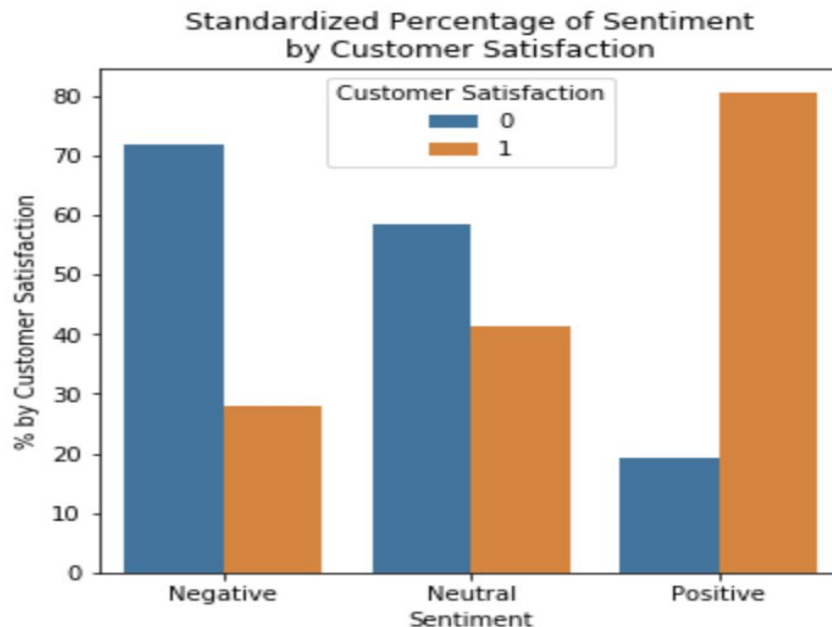
```
Empirical rating mean = 4.183091744741029
95% confidence interval = [4.16833017 4.19741029]
```



Therefore, customers with 'Rating' ≥ 4 were labeled as satisfied customers (1), < 4 were labeled as unsatisfied customers (0).

Exploratory Visualization

This plot below shows that being a satisfied customer is generally a good indicator of positive sentiment, and the same logic applies for the relationship between unsatisfied customers and negative sentiment:



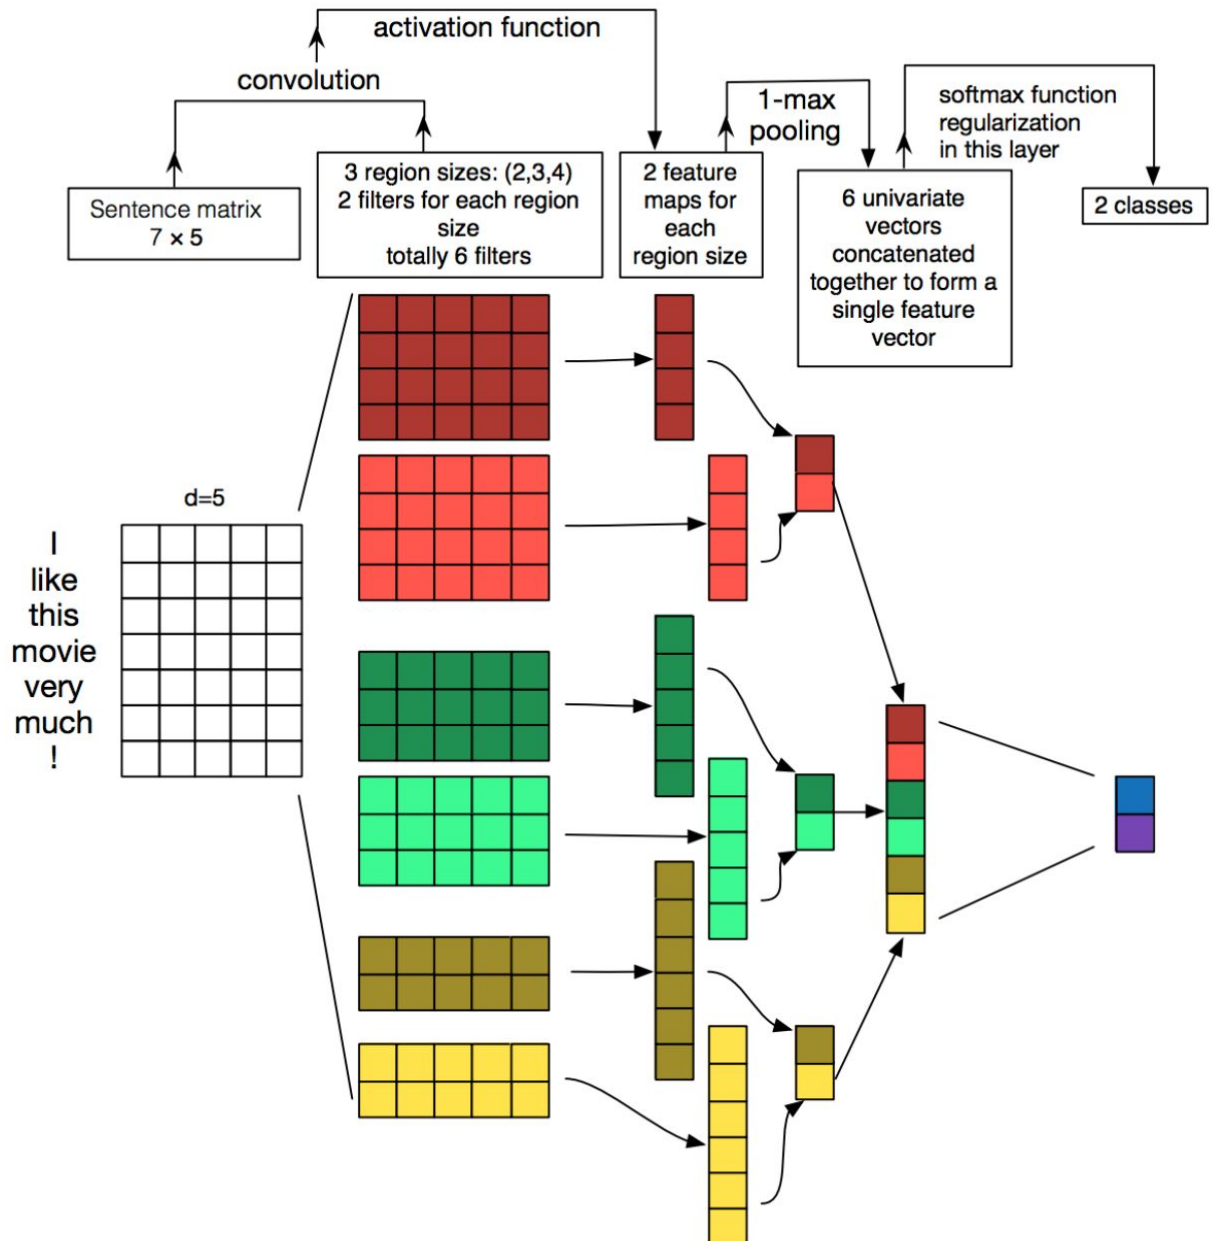
In order to get an idea of what phrases are popular among customers, we created a lookup table to visualize n-grams phrases that are likely to occur among customers. This is helpful for deciding the number of grams when creating CNN channels. Here is the lookup table of the top 10 phrases from satisfied customers:

Satisfied Customers Top Phrases

	2-Gram	Count	4-Gram	Count	6-Gram	Count	8-Gram	Count	10-Gram	Count
0	true size	1203	26 waist 36 hips	29	115 lbs 30 dd 26 waist	10	34b 26 waist 36 hips hem falls inches	3	lightweight soft cotton shorts think meant bea...	2
1	love dress	664	looks great skinny jeans	25	26 waist 36 hips hem falls	7	135 lbs 35 29 36 34d long torso	3	realize ordered like matches look photos short...	2
2	looks great	574	34b 26 waist 36	25	tall 145 lbs 38 32 40	7	wearing medium photos reference measurements 3...	3	outfit hot hot days loose elastic realize orde...	2
3	usually wear	555	love love love dress	23	34b 26 waist 36 hips hem	6	tall 145 lbs 38 36d 32 40 size	3	hot hot days loose elastic realize ordered lik...	2
4	fit perfectly	540	115 lbs 30 dd	21	5ft2in 34b 26 waist 36 hips	5	adorable comfortable 115 lbs ordered xs hits k...	2	hae preferred colors photo big deal guess order...	2
5	fits perfectly	480	dress fits true size	17	regular size small 34d 27 35	5	outfit hot hot days loose elastic realize ordered	2	hot days loose elastic realize ordered like ma...	2
6	size small	427	size small fits perfectly	16	36dd 10 12 tops 12 14	4	weight summer outfit hot hot days loose elastic	2	days loose elastic realize ordered like matche...	2
7	love love	415	lbs 30 dd 26	15	photos reference measurements 38 30 40	4	light weight summer outfit hot hot days loose	2	loose elastic realize ordered like matches loo...	2
8	love top	405	usually wear small medium	14	145 lbs 38 36d 32 40	4	thin light weight summer outfit hot hot days	2	elastic realize ordered like matches look phot...	2
9	usual size	393	great skinny jeans leggings	14	tall 145 lbs 38 36d 32	4	wearing thin light weight summer outfit hot hot	2	ordered like matches look photos shorts low cu...	2

Algorithms and Techniques

The classifier we built was a Multi-Channel Convolutional Network. In this project, we created 3 channels of CNN (different grams), which helped us extract features from an (dropout) embedding layer that was created based on an embedding matrix from a pre-trained Word2Vec model. After a Max-Pooling layer, the network was concatenated, passed through a dropout layer, and ended with a dense layer to perform final interpretation:



Convolutional Neural Network Architecture for Sentence Classification

Taken from "A Sensitivity Analysis of (and Practitioners' Guide to) Convolutional Neural Networks for Sentence Classification", 2015.

The following hyperparameters could be tuned to potentially enhance the performance of the model:

- **Optimizer Hyperparameters**
 - Learning rate
 - Batch size
 - The number of training iterations (Epochs)
 - Optimization Algorithm (Update rule)
 - Dropout rate
 - Weight regularization
- **Model Hyperparameters**
 - Kernel size (n-grams)
 - The number of filters (Feature maps)
 - The number of embedding dimension (Word representations)
 - The number of channels (For carrying different kernels (n-grams))
 - Network depth
 - Layer types/ sequence
 - The use of pre-trained model (Word embeddings)

Benchmark

The model's performance was compared with the top kernel of the dataset on Kaggle. The benchmark model adopts the Multinomial Naive Bayes Algorithm to make predictions. Although the person uses 'Recommended IND' and 'Rating' as target variables, the fact that it is a NLP classification problem does not change. We tested our Neural Network on the same target variables to compare the performance. For 'Recommended IND', the person's model yielded a f-1 score of 0.88. Our goal was to beat its prediction on 'Recommended IND'. Meanwhile, we wanted to see how our model did on detecting customers' sentiments (Emotional tone), which in this setting was a multi-class classification problem (Positive/ Negative/ Neutral).

3. Methodology

Data Preprocessing


The preprocessing done in the 'Text Preprocessing, Model Building & Training' section focused on the 'Review Text' that needed for building a Word2Vec model. Before feeding the text into the Word2Vec class from Gensim, the text was cleaned by going through these steps:

1. Remove stopwords from reviews
2. Remove words that are not alphabetic
3. Remove punctuation from reviews
4. Filter out short words (Minimum character = 2)
5. Tokenize words within each review


The 'cleaned' text was then fed into the Word2Vec class to build a model. The reason why we set out to build this model in the first place was because we wanted to leverage the pre-trained word embedding to build our CNN classifier. A word embedding is a learned representation for text where words that have the same meaning have a similar representation. When developing the word embedding with Word2Vec, here were the parameters defined:

- **size:** 100 (Default) - This is the number of dimensions of the embedding. That is, the length of the dense vector to represent each token (word)
- **window:** 5 (Default) - This is the maximum distance between a target word and words around the target word
- **Min_count:** 1 - This is the minimum count of words to consider when training the model; words with an occurrence less than this count will be ignored
- **workers:** 3 (Default) - The number of threads to use while training
- **sg:** 0 (Default as CBOW) - The training algorithm, either CBOW (0) or skip-gram (1)

The CBOW (Continuous Bag-of-Words) model learns the embedding by predicting the current word based on its context. The continuous skip-gram model learns by predicting the surrounding words given a current word.



After the model was built, it was saved so we could reuse it when building our CNN classifier afterwards. The way we leveraged it was to extract its weights to construct an embedding matrix, which would be used to build the embedding layer in our CNN classifier. Word2Vec is a statistical method for efficiently learning a standalone word embedding from a text corpus, and it could make the neural network based training of the embedding more efficient. Here is a plot of the high dimension word vectors cluster (First 300 words) learn by the model:



Implementation

The implementation process can generally be divided into 2 stages:

1. Encode text data
2. Build CNN model

In the first stage, we first made use of the sklearn function to split the data set into train and test set. We then cleaned the train set and encode the data with the helper functions defined in the scripts of 'nlp.py' and 'cnn.py' respectively. For predicting sentiment types, we needed to numerically encode and categorize the labels beforehand:

1. Numerically encode and categorize labels (For predicting sentiments)
2. Split the data set, with 'Review_Text' as inputs and 'Recommended IND' or 'Sentiments' as labels (Outputs)
3. `clean_reviews()`: apply the helper function from 'nlp.py' to clean the each set

4. `create_tokenizer()`: apply the helper function from 'cnn.py' to create a tokenizer on the train set
5. `max_length()`: apply the helper function from 'cnn.py' to get the maximum review length
6. `encode_reviews()`: apply the helper function from 'cnn.py' to encode and pad all reviews in each set

In the second stage, we defined the model's hyperparameters, built the model with the function from the script 'cnn.py', fit our model and logged the performance, and finally saved the trained model for testing afterwards. In particular, we specify whether or not to use a pre-trained embedding when calling the helper function to build the CNN. Here, we used the pre-trained model mentioned in the Data Preprocessing section:

1. Define model's hyperparameters (Load a pre-trained embedding)
2. `build_cnn()`: apply the helper function from 'cnn.py' to create a multi-channel(s) review classifier
3. Fit the model and start training
4. Save the trained model, and test it on the train set

Here is a summary of our compiled model on predicting customers' sentiments:

Layer (type)	Output Shape	Param #	Connected to
input_3 (InputLayer)	(None, 58)	0	
embedding_3 (Embedding)	(None, 58, 100)	1193500	input_3[0][0]
dropout_5 (Dropout)	(None, 58, 100)	0	embedding_3[0][0]
conv1d_7 (Conv1D)	(None, 58, 150)	30150	dropout_5[0][0]
conv1d_8 (Conv1D)	(None, 58, 150)	60150	dropout_5[0][0]
conv1d_9 (Conv1D)	(None, 58, 150)	75150	dropout_5[0][0]
max_pooling1d_7 (MaxPooling1D)	(None, 57, 150)	0	conv1d_7[0][0]
max_pooling1d_8 (MaxPooling1D)	(None, 57, 150)	0	conv1d_8[0][0]
max_pooling1d_9 (MaxPooling1D)	(None, 57, 150)	0	conv1d_9[0][0]
flatten_7 (Flatten)	(None, 8550)	0	max_pooling1d_7[0][0]
flatten_8 (Flatten)	(None, 8550)	0	max_pooling1d_8[0][0]
flatten_9 (Flatten)	(None, 8550)	0	max_pooling1d_9[0][0]
concatenate_3 (Concatenate)	(None, 25650)	0	flatten_7[0][0] flatten_8[0][0] flatten_9[0][0]
dropout_6 (Dropout)	(None, 25650)	0	concatenate_3[0][0]
activation_3 (Activation)	(None, 25650)	0	dropout_6[0][0]
dense_3 (Dense)	(None, 3)	76953	activation_3[0][0]
Total params: 1,435,903			
Trainable params: 1,435,903			
Non-trainable params: 0			

To recall, since we were developing a system that could potentially predict the sentiment of a textual customer review as either positive, neutral, or negative, the same data preparation was done on both the train set and test set. It is critical that all reviews (In the train set and test set) go through the exact same preprocessing steps, from cleaning to tokenization; so we wrote a script to include all the helper functions needed for text preprocessing and encoding. This ensures objectivity by having all text go through the same ‘assembly line’. On top of that, it helps avoid writing repetitive code.

Refinement

Apart from predicting customers’ sentiments, as mentioned previously, we also wanted to beat the model created in the top kernel, which has an f-1 score of 0.88 on predicting ‘Recommend IND’. Our initial model yielded an accuracy of 0.85 in the validation set over 10 training epochs. To achieve better results, here are the values of parameters we tried to define for our model:

- Kernel (Different filter size/ n-grams): 2, 3, 4, 5, 6
- Feature map (Number of filters): 100, 150, 200, 250, 300
- Dropout: 0.2, 0.3, 0.5
- Batch size: 50, 75, 100, 150, 200

And here are the final settings:

- Kernel (Different filter size/ n-grams): 2, 4, 5
- Feature map (Number of filters): 150
- Dropout: 0.5
- Batch size: 75

When choosing kernel size, we referred to the lookup table created earlier in the project to have a general idea, and we found this approach quite useful. In fact, the kernel value should be tuned for each problem. The same logic applies to feature map. As for dropout, it had little effect on the model’s performance, at least in our case it really did not help much. Our final model yielded an accuracy of 0.88 on predicting ‘Recommended IND’ in the validation. Here is the training log on ‘Recommended IND’ over 10 epochs:

```
Train on 10612 samples, validate on 5227 samples
Epoch 1/10
10612/10612 [=====] - 22s 2ms/step - loss: 2.8566 - acc: 0.8243 - val_loss: 0.5136 - val_ac
c: 0.8303
Epoch 2/10
10612/10612 [=====] - 20s 2ms/step - loss: 0.4336 - acc: 0.8495 - val_loss: 0.4172 - val_ac
c: 0.8506
Epoch 3/10
10612/10612 [=====] - 21s 2ms/step - loss: 0.3805 - acc: 0.8602 - val_loss: 0.3820 - val_ac
c: 0.8619
Epoch 4/10
10612/10612 [=====] - 20s 2ms/step - loss: 0.3676 - acc: 0.8668 - val_loss: 0.4058 - val_ac
c: 0.8445
Epoch 5/10
10612/10612 [=====] - 21s 2ms/step - loss: 0.3416 - acc: 0.8774 - val_loss: 0.3547 - val_ac
c: 0.8638
Epoch 6/10
10612/10612 [=====] - 20s 2ms/step - loss: 0.3262 - acc: 0.8832 - val_loss: 0.3329 - val_ac
c: 0.8802
Epoch 7/10
10612/10612 [=====] - 20s 2ms/step - loss: 0.3020 - acc: 0.8928 - val_loss: 0.3345 - val_ac
c: 0.8797
Epoch 8/10
10612/10612 [=====] - 22s 2ms/step - loss: 0.2959 - acc: 0.8960 - val_loss: 0.3285 - val_ac
c: 0.8846
Epoch 9/10
10612/10612 [=====] - 20s 2ms/step - loss: 0.2833 - acc: 0.9023 - val_loss: 0.3237 - val_ac
c: 0.8888
Epoch 10/10
10612/10612 [=====] - 21s 2ms/step - loss: 0.2680 - acc: 0.9114 - val_loss: 0.3342 - val_ac
c: 0.8822
```

4. Results

Model Evaluation and Validation

The final model architecture and hyperparameters, as illustrated in the ‘Algorithms and Techniques’ section, were chosen based on its performance compared with the benchmark model. The predictions on ‘Sentiments’ were actually better than what we expected. For both types of predictions, we generally used the same set of hyperparameters, except for activation function, loss function, and output units, as the predictions on ‘Sentiments’ is a multi-class problem. Here is the list of our definitions on hyperparameters:

- Embedding dimension: 100
- The number of CNN channels: 3
- Filter size (ngrams): 2, 4, 5
- Feature map size: 150
- Drop out rate: 0.5
- Weight regularization: L2 (0.03)
- Optimizer: Adam
- Batch size: 75
- The number of training epochs: 10
- Pre-trained embedding: Word2Vec on all reviews

To verify the model’s robustness, the model after each training epoch was test again a validation set to make sure loss and accuracy generally had the tendency to improve. Finally, the model was tested on the test set, which consisted of 6789 observations that the model had never seen before. Here are training log and the classification summary on ‘Recommended IND’ and ‘Sentiments’ respectively:

```
Train on 10612 samples, validate on 5227 samples
Epoch 1/10
10612/10612 [=====] - 22s 2ms/step - loss: 2.8566 - acc: 0.8243 - val_loss: 0.5136 - val_ac
c: 0.8303
Epoch 2/10
10612/10612 [=====] - 20s 2ms/step - loss: 0.4336 - acc: 0.8495 - val_loss: 0.4172 - val_ac
c: 0.8506
Epoch 3/10
10612/10612 [=====] - 21s 2ms/step - loss: 0.3805 - acc: 0.8602 - val_loss: 0.3820 - val_ac
c: 0.8619
Epoch 4/10
10612/10612 [=====] - 20s 2ms/step - loss: 0.3676 - acc: 0.8668 - val_loss: 0.4058 - val_ac
c: 0.8445
Epoch 5/10
10612/10612 [=====] - 21s 2ms/step - loss: 0.3416 - acc: 0.8774 - val_loss: 0.3547 - val_ac
c: 0.8638
Epoch 6/10
10612/10612 [=====] - 20s 2ms/step - loss: 0.3262 - acc: 0.8832 - val_loss: 0.3329 - val_ac
c: 0.8802
Epoch 7/10
10612/10612 [=====] - 20s 2ms/step - loss: 0.3020 - acc: 0.8928 - val_loss: 0.3345 - val_ac
c: 0.8797
Epoch 8/10
10612/10612 [=====] - 22s 2ms/step - loss: 0.2959 - acc: 0.8960 - val_loss: 0.3285 - val_ac
c: 0.8846
Epoch 9/10
10612/10612 [=====] - 20s 2ms/step - loss: 0.2833 - acc: 0.9023 - val_loss: 0.3237 - val_ac
c: 0.8888
Epoch 10/10
10612/10612 [=====] - 21s 2ms/step - loss: 0.2680 - acc: 0.9114 - val_loss: 0.3342 - val_ac
c: 0.8822
```

	precision	recall	f1-score	support
Not Recommended	0.75	0.62	0.67	1167
Recommended	0.92	0.96	0.94	5622
avg / total	0.89	0.90	0.89	6789

Train on 10612 samples, validate on 5227 samples

```
Epoch 1/10
10612/10612 [=====] - 21s 2ms/step - loss: 2.6905 - acc: 0.9208 - val_loss: 0.3328 - val_ac
c: 0.9327
Epoch 2/10
10612/10612 [=====] - 20s 2ms/step - loss: 0.2933 - acc: 0.9287 - val_loss: 0.2663 - val_ac
c: 0.9327
Epoch 3/10
10612/10612 [=====] - 21s 2ms/step - loss: 0.2624 - acc: 0.9303 - val_loss: 0.2601 - val_ac
c: 0.9286
Epoch 4/10
10612/10612 [=====] - 24s 2ms/step - loss: 0.2500 - acc: 0.9307 - val_loss: 0.2500 - val_ac
c: 0.9288
Epoch 5/10
10612/10612 [=====] - 22s 2ms/step - loss: 0.2345 - acc: 0.9313 - val_loss: 0.2676 - val_ac
c: 0.9151
Epoch 6/10
10612/10612 [=====] - 22s 2ms/step - loss: 0.2263 - acc: 0.9324 - val_loss: 0.2310 - val_ac
c: 0.9348
Epoch 7/10
10612/10612 [=====] - 21s 2ms/step - loss: 0.2100 - acc: 0.9367 - val_loss: 0.2276 - val_ac
c: 0.9334
Epoch 8/10
10612/10612 [=====] - 19s 2ms/step - loss: 0.2003 - acc: 0.9373 - val_loss: 0.2310 - val_ac
c: 0.9296
Epoch 9/10
10612/10612 [=====] - 20s 2ms/step - loss: 0.1930 - acc: 0.9411 - val_loss: 0.2288 - val_ac
c: 0.9309
Epoch 10/10
10612/10612 [=====] - 20s 2ms/step - loss: 0.1856 - acc: 0.9450 - val_loss: 0.2339 - val_ac
c: 0.9330
```

	precision	recall	f1-score	support
Positive	0.94	0.99	0.97	6315
Negative	0.63	0.17	0.27	436
Neutral	0.00	0.00	0.00	38
avg / total	0.92	0.93	0.92	6789

Justification

As shown in the previous section, our model yielded a f-1 score of 0.89 and 0.92 on classifying ‘Recommended IND’ and ‘Sentiments’ respectively. For ‘Recommended IND’, we successfully beat the benchmark model by a slightly amount of 0.01. For ‘Sentiments’, the classification performance was even better; so we were quite successful in the scope of this project.

To caveat, the results were based on a relatively small data set. To recall, we only had 22628 observations after cleaning the data set. And the word embedding was trained on about 14000 vocabularies. In reality, embeddings are learned from much larger corpora of text, which sometimes could be billions of words. In a nutshell, if we wanted to deploy the model on the

e-commerce site for classifying customers' reviews, we'd probably need more data and let our model 'study' more text.

5. Conclusion

Free Form Visualization

We wanted to see how the model performed without the presence of a pre-trained embedding; so we went ahead and trained the model with the same set of hyperparameters on classifying 'Sentiments' without a pre-trained (Word2Vec) embedding. Here are the training log and the model's final performance on the test set:

```
Train on 10612 samples, validate on 5227 samples
Epoch 1/10
10612/10612 [=====] - 22s 2ms/step - loss: 2.1576 - acc: 0.9233 - val_loss: 0.2759 - val_ac
c: 0.9327
Epoch 2/10
10612/10612 [=====] - 23s 2ms/step - loss: 0.2385 - acc: 0.9297 - val_loss: 0.2540 - val_ac
c: 0.9350
Epoch 3/10
10612/10612 [=====] - 21s 2ms/step - loss: 0.1975 - acc: 0.9404 - val_loss: 0.2628 - val_ac
c: 0.9338
Epoch 4/10
10612/10612 [=====] - 25s 2ms/step - loss: 0.1737 - acc: 0.9509 - val_loss: 0.2564 - val_ac
c: 0.9348
Epoch 5/10
10612/10612 [=====] - 23s 2ms/step - loss: 0.1487 - acc: 0.9600 - val_loss: 0.2708 - val_ac
c: 0.9313
Epoch 6/10
10612/10612 [=====] - 22s 2ms/step - loss: 0.1326 - acc: 0.9650 - val_loss: 0.2780 - val_ac
c: 0.9196
Epoch 7/10
10612/10612 [=====] - 24s 2ms/step - loss: 0.1190 - acc: 0.9703 - val_loss: 0.2913 - val_ac
c: 0.9263
Epoch 8/10
10612/10612 [=====] - 24s 2ms/step - loss: 0.1041 - acc: 0.9759 - val_loss: 0.3242 - val_ac
c: 0.9313
Epoch 9/10
10612/10612 [=====] - 23s 2ms/step - loss: 0.1041 - acc: 0.9750 - val_loss: 0.3922 - val_ac
c: 0.9290
Epoch 10/10
10612/10612 [=====] - 21s 2ms/step - loss: 0.0961 - acc: 0.9781 - val_loss: 0.3546 - val_ac
c: 0.9208
```

	precision	recall	f1-score	support
Positive	0.95	0.97	0.96	6315
Negative	0.38	0.25	0.30	436
Neutral	0.00	0.00	0.00	38
avg / total	0.90	0.92	0.91	6789

As shown above, there was not much of a difference regarding the model's final performance on the test set. One thing worth mentioning though, is the model's performance in the validation set. The model without pre-trained embedding seemed to perform worse as training epoch increased.

Reflection


The entire process used in this project could be summarized by the following:

1. Discover a problem, and define the scope of the project (How to approach the problem)
2. Data gathering
3. Setting a benchmark that the model can be compared with

4. Preliminary investigation of the data set
5. Text preprocessing and model building
6. Model training and refinement
7. Model testing and refinement

As with most projects we've done before, interesting insights usually start to show up early in the process. Like in the preliminary investigation phase in this project, we got to discover a 'Rating' of 4 is a threshold between satisfied and unsatisfied customers, which in turn is actually a very good indicator of customers' sentiments. We went on to check the recommended rate of these customers, and found that the rates (Conversions) are 0.99 and 0.25 for satisfied and unsatisfied customers respectively. The difference was even proved significant based on permutation (10000 simulations), which helped us reject the null hypothesis that there is no difference in recommended rate between satisfied and unsatisfied customers:

Empirical Recommended Rate Difference = 0.743



Again, with most Machine Learning problems, hyperparameters tuning is as much a science as an art. That is why we found steps 5-7 the most challenging. It took patience and trial-and-error to yield good results because at the end of the day, how to define the model and values really depends on what sort of problem we are dealing with. Having that said, this is what makes Machine Learning interesting!

Improvement

As mentioned previously, one possible way to enhance the model's performance is to train it on a larger and more representative data set. This would mean a data set that includes more observations and a much larger corpora of text (For embeddings). Regarding model architecture and complexity, here are a couple of aspects that we could think of:

- N-grams: This is the kernel size/ filter size in the CNN. Different problems supposedly require different combinations of n-grams
- The number of CNN channels: The number of channels could have great impact on performance
- Network depth: Although not necessary, deeper model usually perform better

If we were to extend the project, we'd like to try deeper network or perhaps construct CNNs on character level. This would mean greater computational power and labor-intensive effort (For text preprocessing) are required.

6. Reference

- [Twitter Sentiment Classification using Distant Supervision](#)
- [Convolutional Neural Networks for Sentence Classification](#)
- [Convolutional Neural Networks for Sentence Classification \(code\)](#)
- [Best Practices for Document Classification with Deep Learning](#)
- [Character-level Convolutional Networks for Text Classification](#)
- [A Sensitivity Analysis of \(and Practitioners' Guide to\) Convolutional Neural Networks for Sentence Classification](#)
- [Natural Language Processing \(almost\) from Scratch](#)
- [A Primer on Neural Network Models for Natural Language Processing](#)
- [A Sentimental Education: Sentiment Analysis Using Subjectivity Summarization Based on Minimum Cuts](#)
- [Understanding Convolutional Neural Networks for NLP](#)
- [Guided Numeric and Text Exploration E-Commerce](#)
- [List of NLP resources](#)
- [Word Cloud with Python](#)
- [Deep Learning for NLP Best Practices](#)
- [Embed, encode, attend, predict: The new deep learning formula for state-of-the-art NLP models](#)